

2024. LegalBench-RAG: A Domain-Specific Benchmark for Evaluating Retrieval in Legal RAG Systems. *arXiv preprint* (2024). arXiv:2408.10343 <https://arxiv.org/abs/2408.10343>
- [31] Yang Wang and Hassan A Karimi. 2024. Exploring Large Language Models for Climate Forecasting. *arXiv preprint arXiv:2411.13724* (2024).
- [32] ZhiLin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W Cohen, Ruslan Salakhutdinov, and Christopher D Manning. 2018. HotpotQA: A dataset for diverse, explainable multi-hop question answering. *arXiv preprint arXiv:1809.09600* (2018).
- [33] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. React: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*.
- [34] Siyu Yuan, Jiangjie Chen, Ziquan Fu, Xuyang Ge, Soham Shah, Charles Jankowski, Yanghua Xiao, and Deqing Yang. 2023. Distilling Script Knowledge from Large Language Models for Constrained Language Planning. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 4303–4325.
- [35] Zhiyuan Zeng, Jiashuo Liu, Siyuan Chen, Tianci He, Yali Liao, Yixiao Tian, Jinpeng Wang, Zaiyuan Wang, Yang Yang, Lingyue Yin, et al. 2025. FutureX: An Advanced Live Benchmark for LLM Agents in Future Prediction. *arXiv preprint arXiv:2508.11987* (2025).
- [36] Yuxuan Zhang, Zhiyuan Zhang, Yicheng Wang, Yuxuan Su, Yixuan Su, and Yixuan Su. 2023. PubHealth: A Benchmark for Public Health Question Answering. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, Toronto, Canada, 9802–9822. doi:10.18653/v1/2023.acl-long.546
- [37] Siyun Zhao, Yuqing Yang, Zilong Wang, Zhiyuan He, Luna K Qiu, and Lili Qiu. 2024. Retrieval augmented generation (rag) and beyond: A comprehensive survey on how to make your llms use external data more wisely. *arXiv preprint arXiv:2409.14924* (2024).
- [38] Chang Zong, Yuchen Yan, Weiming Lu, Jian Shao, Eliot Huang, Heng Chang, and Yueting Zhuang. 2024. Triad: A Framework Leveraging a Multi-Role LLM-based Agent to Solve Knowledge Base Question Answering. arXiv:2402.14320 [cs.CL] <https://arxiv.org/abs/2402.14320>

Appendix

Case Studies

To evaluate the proposed supportiveness metric in RAG, we consider two representative cases shown in Figure 6. The first question asks whether the total crypto market cap will exceed \$4 trillion on January 20, 2025. Among the retrieved articles, the blue-highlighted news discusses substantial rebounds in Bitcoin and Ether alongside ETF approvals—information directly indicative of a “Yes” outcome. Conversely, the white-highlighted news describes global market declines and cryptocurrency crashes, contradicting the supportive stance. Our metric assigns higher scores to the supportive article, reflecting its stronger evidential link to the answer.

The second question concerns whether Joe Rogan will attend the presidential inauguration. The supportive article details Rogan’s past political activities and his collaboration with Donald Trump during the campaign, which strengthens the likelihood of his attendance. In contrast, the non-supportive article mentions Rogan only in the context of broader shifts in media influence, without providing any concrete statement or event suggesting his presence at the inauguration. Our metric correctly ranks the supportive article higher, aligning with the ground truth answer “Yes”.

Across both cases, the metric effectively distinguishes articles offering direct, answer-relevant evidence from those that are tangential or context-only, thereby enhancing RAG answer accuracy.

Construction Details

During constructing, we use gpt-4o-mini-2024-07-18 for all LLM callings. We set window size w to 3 which is enough large in our pilot study. For computing each probability in CIL, we call twice gpt-4o-mini-2024-07-18 and get the average score. The constructing prompts we use are shown in prompts (a-f).

Prompts

We list all prompts in the following Figures (a-h).

Agentic RAG Tool

We show the article retrieval tool for all agentic RAG methods:

```

1 class AgenticRAGTool:
2     name = "query_rag"
3     description = (
4         "To search in a previously built RAG index
5             to find the most relevant chunks of
6                 text."
7     )
8     parameters = [
9         {
10             'name': 'query',
11             'type': 'string',
12             'description': 'the query text to
13                 search for relevant text chunks.',
14             'required': True
15         },
16         {
17             'name': 'top_k',
18             'type': 'integer',
19             'description': 'the number of top
20                 relevant chunks to retrieve (
21                     default is 3).',
22             'required': False
23         }
24     ]

```

Listing 1: Python Class Definition for Agentic RAG Tool

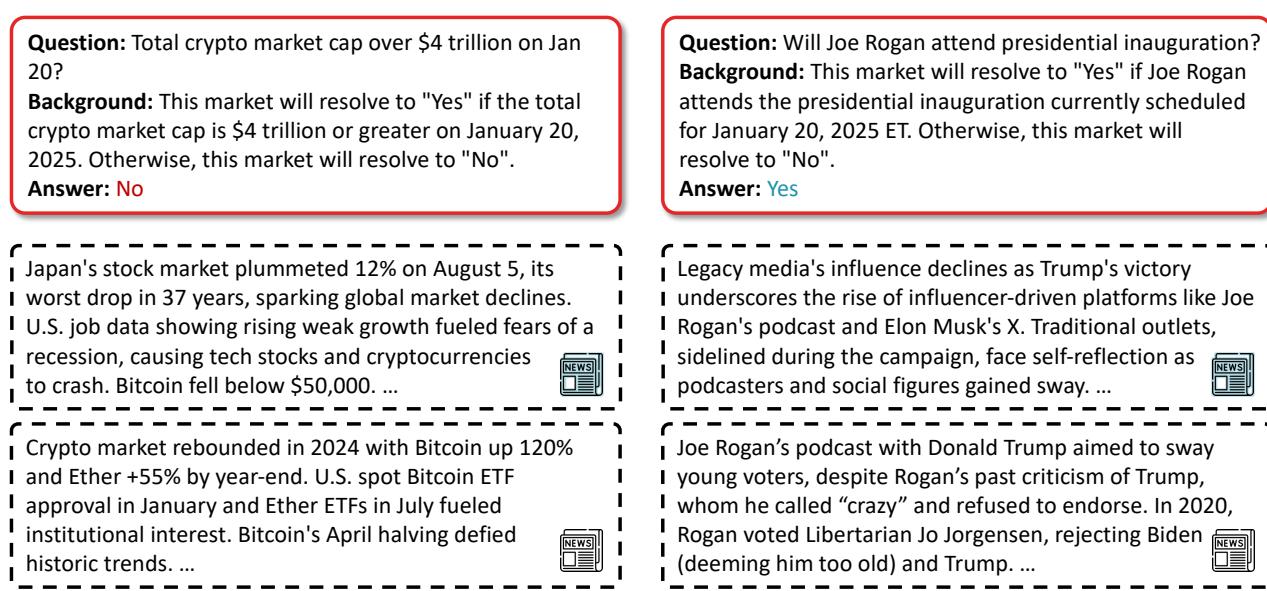


Figure 6: Case studies. The supportive and non-supportive articles are selected by CIL scores.

(a) Entity Query Generation

I will provide you with a forecasting question and the background information for the question. Extract the named entities, events of the question. Each entity and event are up to 5 words. The named entities can only be people, organizations, countries, locations while can not be date or time. Put all result items in a list that I can parse by JSON as ["entity 1", "entity 2", "event 1", "event 2", ...].

Question: *Q*

Question Background: *B*

Question Date: *D*

Output:

(b) Resolving Steps Query Generation

I will provide you with a forecasting question and the background information for the question. I will then ask you to generate short search queries (up to max words words each) that I'll use to find articles on Google News to help answer the question. The articles should be mainly about event arguments such as subjects, objects, locations, organizations of the events in question and background information. You must generate this exact amount of queries: num keywords. Put all result items in a list that I can parse by JSON as ["step 1", "step 2", "step 3", ...].

Question: *Q*

Question Background: *B*

Question Date: *D*

Output:

(c) Similar Events Query Generation

I will provide you with a forecasting question and the background information for the question. I will then ask you to generate short search queries (up to max words words each) that I'll use to find articles of similar events on Google News to help answer the question. The similar events are events happened on other similar entities in the history. Or events happened on question entities but on other date. You must generate this exact amount of queries: num keywords. Put all result items in a list that I can parse by JSON as ["event 1", "event 2", "event 3", ...].

Question: *Q*

Question Background: *B*

Question Date: *D*

Output: